

Python &Deep Learning Lab-3

Team ID-06 Team

Members

1.Tarun Teja Kasturi -11

2.Lakshmana Kumar Mettu -16

3.Pavan Kumar Chongala -04

- Keras
 - Tensor board
 - Google Colaboratory
 - Pycharm
-

Task-1:Build a Sequential model using keras to implement Linear Regression with heart uci dataset

```

#Linear Regression on Heart UCI data set
#importing the required libraries
import pandas as pd

from keras.models import Sequential
from keras.layers import Dense
from keras import optimizers
from keras.datasets import mnist
from keras.utils import np_utils
from keras.callbacks import TensorBoard
import tensorflow as tf
import matplotlib.pyplot as plt

#setting the batch size and epochs
from sklearn.model_selection import train_test_split

batch_size = 128
nb_classes = 10
nb_epoch = 30

#loading the dataset
dataset = pd.read_csv("C:/Users/laksh/PycharmProjects/lab-3-DL/LOGISTIC/heart.csv", header=None).values
#print(dataset)
import numpy as np
X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:13], dataset[:,13],
                                                    test_size=0.25, random_state=87)

##Data normalization
X_train = X_train.astype(np.float)
X_test = X_test.astype(np.float)
X_train /= 255
X_test /= 255
print(X_train)
Y_Train = np_utils.to_categorical(Y_train, nb_classes)
Y_Test = np_utils.to_categorical(Y_test, nb_classes)

```



```

# Linear regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='relu'))
model.compile(optimizer='rmsprop', loss='mean_absolute_error', metrics=['accuracy'])
model.summary()

#Tensorboard log generation for graphs
tensorboard = TensorBoard(log_dir="logs/{}", histogram_freq=0, write_graph=True, write_images=True)

#fitting the model
history=model.fit(X_train, Y_Train, nb_epoch=nb_epoch, batch_size=batch_size,callbacks=[tensorboard])

#predicting the accuracy of the model
score = model.evaluate(X_test, Y_Test, verbose=1)
print('Loss: %.2f, Accuracy: %.2f' % (score[0], score[1]))

#plotting the loss
plt.plot(history.history['loss'])
# plt.plot(history.history['test_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()


```

```
C:\Users\laksh\AppData\Local\Programs\Python\Python36\python.exe C:/Users/laksh/PycharmProjects/lab-3-DL/que-1.py
Using TensorFlow backend.
[[0.22745098 0.00392157 0. ... 0.00784314 0. ... 0.01176471]
 [0.23529412 0. ... 0.01176471 ... 0.00784314 0. ... 0.00784314]
 [0.24313725 0. ... 0. ... 0.00392157 0.01176471 0.00784314]
 ...
 [0.16862745 0.00392157 0. ... 0.00392157 0.01568627 0.01176471]
 [0.2627451 0. ... 0. ... 0.00784314 0.00784314 0.00784314]
 [0.21176471 0.00392157 0. ... 0.00392157 0.00392157 0.01176471]]
C:/Users/laksh/PycharmProjects/lab-3-DL/que-1.py:37: UserWarning: Update your 'Dense' call to the Keras 2 API: 'Dense(input_shape=(13,), activation='relu', units=10, kernel_initializer='normal'
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='relu'))

Layer (type) Output Shape Param #
-----
dense_1 (Dense) (None, 10) 140
-----
Total params: 140
Trainable params: 140
Non-trainable params: 0
```

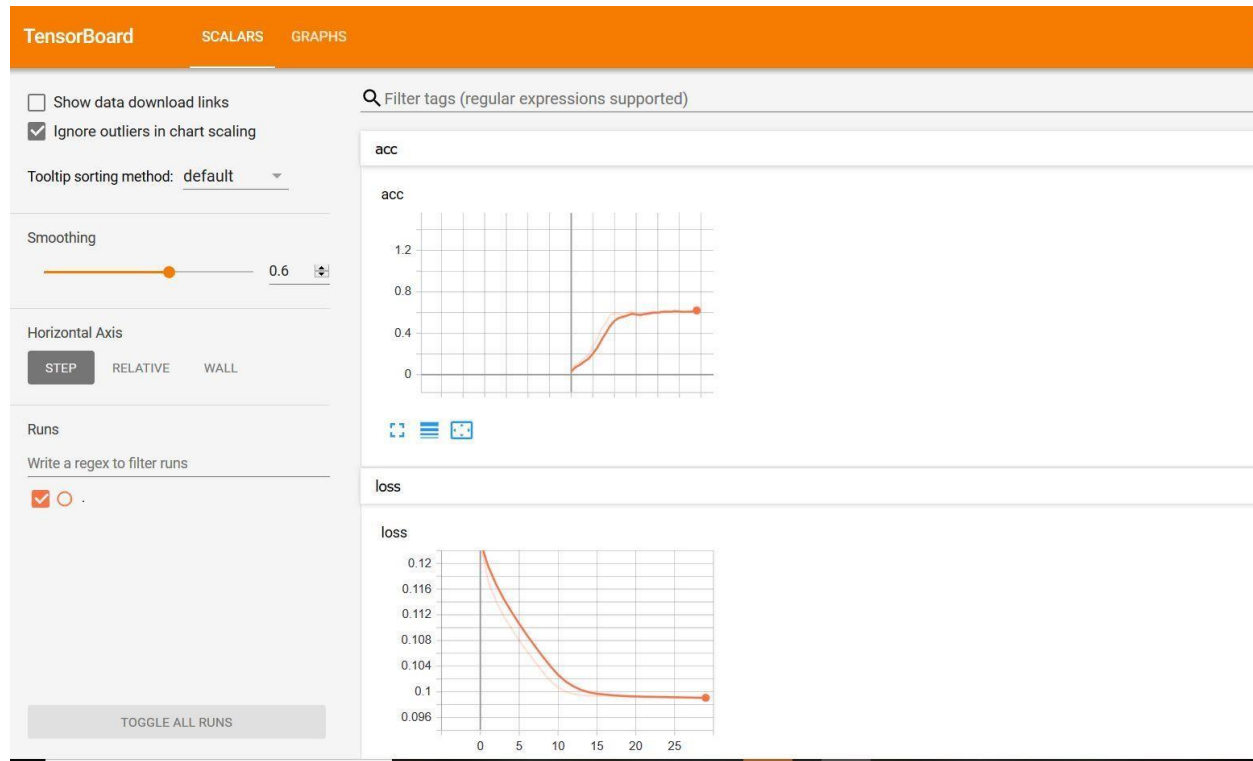
```
32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 197us/step
Loss: 0.10, Accuracy: 0.64
```

Figure 1



Tensor board commands and loss and accuracy curves plotted on tensor board:

```
C:\Users\laksh\PycharmProjects\lab-3-DL>tensorboard --logdir="C:\Users\laksh\PycharmProjects\lab-3-DL\logs\{"  
TensorBoard 1.13.1 at http://DESKTOP-CFC46NE:6006 (Press CTRL+C to quit)
```



Changing hyper parameters such as batch size , no.of.epochs,learning rate and activation type and optimizer and plotting the loss ,observing the accuracy and loss along with tensor board graph for loss,accuracy are given below.

Highest accuracy achieved as 0.64 for epoch -30

Optimizer -RMS prop

Activation type-Relu

```

batch_size = 64
nb_classes = 10
nb_epoch = 10

#loading the dataset
dataset = pd.read_csv("C:/Users/laksh/PycharmProjects/lab-3-DL/LOGISTIC/heart.csv", header=None).values
#print(dataset)
import numpy as np
X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:13], dataset[:,13],
                                                    test_size=0.25, random_state=87)

##Data normalization
X_train = X_train.astype(np.float)
X_test = X_test.astype(np.float)
X_train /= 255
X_test /= 255
print(X_train)
Y_train = np_utils.to_categorical(Y_train, nb_classes)
Y_test = np_utils.to_categorical(Y_test, nb_classes)

# Linear regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='relu'))
model.compile(optimizer='rmsprop', loss='mean_absolute_error', metrics=['accuracy'])
model.summary()

#tensorboard log generation for graphs


```

```

32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 184us/step
Loss: 0.10, Accuracy: 0.58

```

Figure 1



```

batch_size = 128
nb_classes = 10
nb_epoch = 5

#loading the dataset
dataset = pd.read_csv("C:/Users/laksh/PycharmProjects/lab-3-DL/LOGISTIC/heart.csv", header=None).values
#print(dataset)
import numpy as np
X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:13], dataset[:,13],
                                                    test_size=0.25, random_state=87)

##Data normalization
X_train = X_train.astype(np.float)
X_test = X_test.astype(np.float)
X_train /= 255
X_test /= 255
print(X_train)
Y_train = np_utils.to_categorical(Y_train, nb_classes)
Y_test = np_utils.to_categorical(Y_test, nb_classes)

# L regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='relu'))
model.compile(optimizer='rmsprop', loss='mean_absolute_error', metrics=['accuracy'])
model.summary()


#Tensorboard log generation for graphs

```

```

32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 198us/step
Loss: 0.10, Accuracy: 0.64
|

```



```

# Linear regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='sigmoid'))
model.compile(optimizer='SGD', loss='mean_absolute_error', metrics=['accuracy'])
model.summary()

```

```

32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 236us/step
Loss: 0.50, Accuracy: 0.49
|

```

Task2:Implemeting Logistic Regression on on Heart Disease UCI Dataset

Code snippets and output observations by varying hyper parameters as well as optimization and activation types are given below. Here also we have changed the epochs ,batch size ,learning rate as well as Activation type and optimizers such as SGD,RMS prop.

Maximum accuracy-0.70 for softmax(Activation type) ,SGD(Optimizer)

```

# importing libraries
import pandas as pd
import matplotlib.pyplot as plt
from keras.models import Sequential
from keras.layers import Dense
from keras import optimizers
from keras.datasets import mnist
from keras.utils import np_utils
from keras.callbacks import TensorBoard
import tensorflow as tf

# setting batch and epochs
from sklearn.model_selection import train_test_split

batch_size = 140
nb_classes = 10
nb_epoch = 100

# loading the dataset(mnist)
dataset = pd.read_csv("C:/Users/laksh/PycharmProjects/lab-3-DL/LOGISTIC/heart.csv", header=None).values
#print(dataset)
import numpy as np
X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:13], dataset[:,13],
                                                    test_size=0.25, random_state=87)
#(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.reshape(100, 784)
X_test = X_test.reshape(5000, 784)
print(X_train)
X_train = X_train.astype(np.float)
X_test = X_test.astype(np.float)
X_train /= 255
X_test /= 255
Y_train = np_utils.to_categorical(Y_train, nb_classes)
Y_test = np_utils.to_categorical(Y_test, nb_classes)

```



```

#performing Logistic regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='softmax'))
model.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

#tensorboard graph generation
tensorboard = TensorBoard(log_dir="logs101/{}", histogram_freq=0, write_graph=True, write_images=True)
history=model.fit(X_train, Y_Train, nb_epoch=nb_epoch, batch_size=batch_size,callbacks=[tensorboard])

#predicting the accuracy of the model
score = model.evaluate(X_test, Y_Test, verbose=1)
print('Loss: %.2f, Accuracy: %.2f' % (score[0], score[1]))

#plotting the loss
plt.plot(history.history['loss'])
# plt.plot(history.history['test_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```

```

16/16 [-----] - 0s 209us/step

```

```


Loss: 1.17, Accuracy: 0.64

```

```

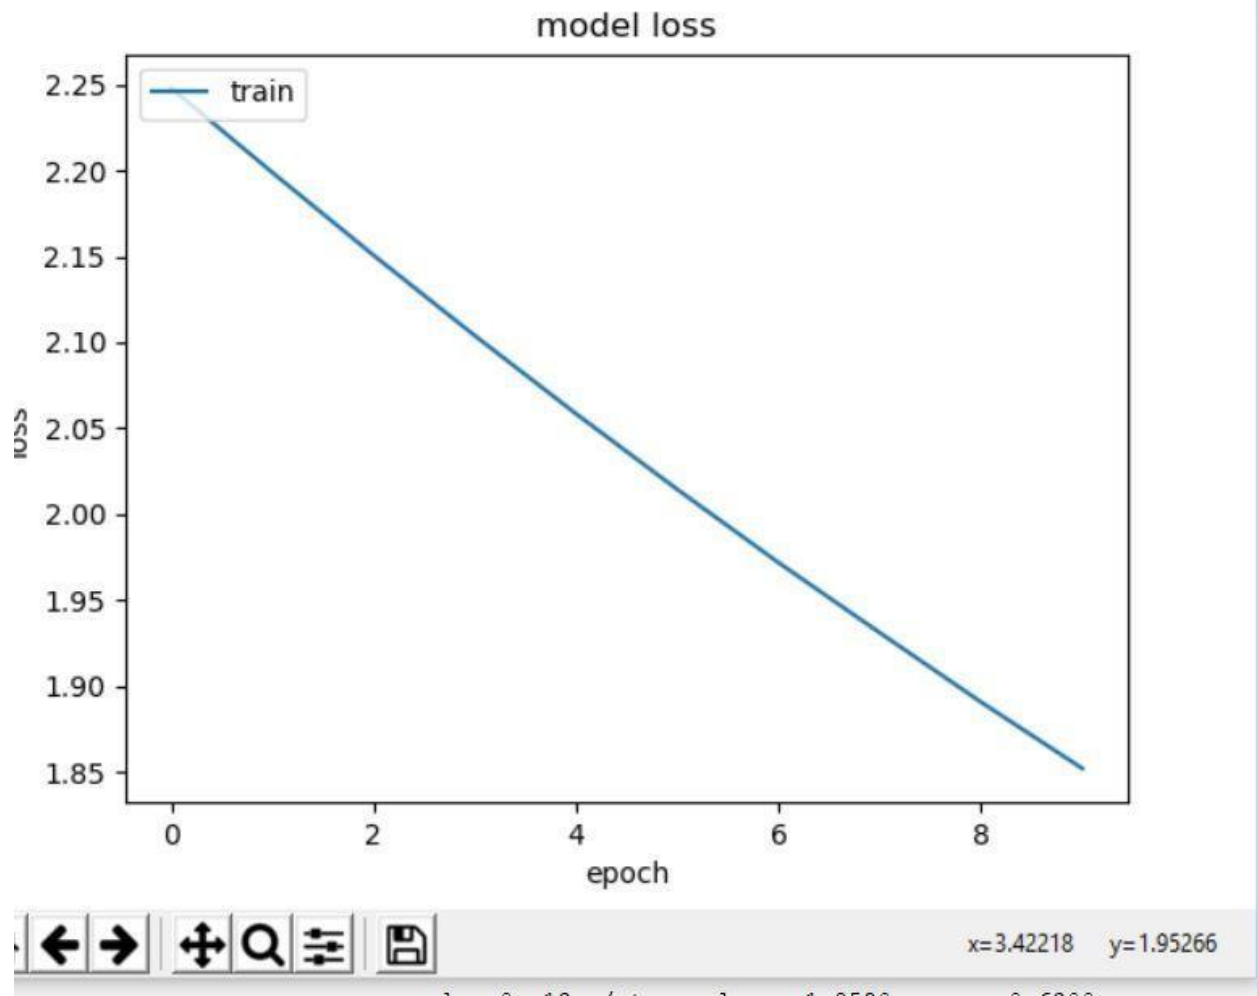
|

```





```
batch_size = 50
nb_classes = 10
nb_epoch = 10
```

```
32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 209us/step
Loss: 2.01, Accuracy: 0.36
```



```
#performing Logistic regression
model = Sequential()
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='softmax'))
model.compile(optimizer='RMSprop', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```



x=5.7375 y=2.1654

```
C:\Users\laksh\AppData\Local\Programs\Python\python36\python.exe C:/Users/laksh/PycharmProjects/lab-3-DL/task-2.py
Using TensorFlow backend.
[[59.  1.  0. ...  2.  0.  3.]
 [60.  0.  3. ...  2.  0.  2.]
 [62.  0.  0. ...  1.  3.  2.]
 ...
 [43.  1.  0. ...  1.  4.  3.]
 [67.  0.  0. ...  2.  2.  2.]
 [54.  1.  0. ...  1.  1.  3.]]
C:/Users/laksh/PycharmProjects/lab-3-DL/task-2.py:38: UserWarning: Update your 'Dense' call to the Keras 2 API: 'Dense(input_shape=(13,), activation='softmax', units=10, kernel_initializer='normal')'
model.add(Dense(output_dim=10, input_shape=(13,), init='normal', activation='softmax'))

Layer (type)                Output Shape              Param #
-----
dense_1 (Dense)              (None, 10)                140
-----
Total params: 140
Trainable params: 140
Non-trainable params: 0

C:/Users/laksh/PycharmProjects/lab-3-DL/task-2.py:44: UserWarning: The 'nb_epoch' argument in 'fit' has been renamed 'epochs'.
history=model.fit(X_train, Y_train, nb_epoch=nb_epoch, batch_size=batch_size, callbacks=[tensorboard])
2019-04-29 00:28:10.114889: I tensorflow/core/platform/cpu_feature_guard.cc:141] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX2
Epoch 1/5

10/227 [>.....] - ETA: 1s - loss: 2.2328 - acc: 0.1000
227/227 [=====] - 0s 272us/step - loss: 2.1399 - acc: 0.4602
Epoch 2/5

10/227 [>.....] - ETA: 0s - loss: 2.0649 - acc: 0.6000
227/227 [=====] - 0s 61us/step - loss: 2.0204 - acc: 0.4890
Epoch 3/5
```

```

epoch 1/5

10/227 [>.....] - ETA: 1s - loss: 2.2328 - acc: 0.1000
227/227 [=====] - 0s 272us/step - loss: 2.1399 - acc: 0.4802
Epoch 2/5

10/227 [>.....] - ETA: 0s - loss: 2.0649 - acc: 0.6000
227/227 [=====] - 0s 61us/step - loss: 2.0204 - acc: 0.4890
Epoch 3/5

10/227 [>.....] - ETA: 0s - loss: 1.9703 - acc: 0.7000
227/227 [=====] - 0s 52us/step - loss: 1.9163 - acc: 0.4890
Epoch 4/5

10/227 [>.....] - ETA: 0s - loss: 1.8692 - acc: 0.4000
227/227 [=====] - 0s 52us/step - loss: 1.8175 - acc: 0.4934
Epoch 5/5

10/227 [>.....] - ETA: 0s - loss: 1.7912 - acc: 0.3000
227/227 [=====] - 0s 51us/step - loss: 1.7251 - acc: 0.4890

32/76 [=====>.....] - ETA: 0s
76/76 [=====] - 0s 210us/step
Loss: 1.67, Accuracy: 0.36

```

Tensor Board Graphs for loss and accuracy:



Task-3:Image classification using CNN on 10 monkey species dataset

We had completed the classification for given data and labeled with their names. Corresponding snippets are given below. We got validation accuracy as 98.16%.

```
[1] import os
import cv2
import glob
import h5py
import shutil
import imgaug as aug
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
import imgaug.augmenters as iaa
from os import listdir, makedirs, getcwd, remove
from os.path import isfile, join, abspath, exists, isdir, expanduser
from pathlib import Path
from skimage.io import imread
from skimage.transform import resize
from keras.models import Sequential, Model, load_model
from keras.applications.vgg16 import VGG16, preprocess_input
from keras.layers import Conv2D, MaxPooling2D, Dense, Dropout, Input, Flatten
from keras.optimizers import Adam, SGD, RMSprop
from keras.callbacks import ModelCheckpoint, Callback, EarlyStopping
from keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from mlxtend.plotting import plot_confusion_matrix
from sklearn.metrics import confusion_matrix
from mlxtend.plotting import plot_confusion_matrix
from keras import backend as K
import tensorflow as tf

color = sns.color_palette()
%matplotlib inline
%config InlineBackend.figure_format="svg"
```

```
[2] # Set the seed for hash based operations in python
os.environ['PYTHONHASHSEED'] = '0'

seed=1234

# Set the numpy seed
np.random.seed(seed)

# Set the random seed in tensorflow at graph level
#tf.set_random_seed(seed)

# Make the augmentation sequence deterministic
aug.seed(seed)
```

```
[3] from google.colab import drive
drive.mount('/content/drive')
```

📁 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[4] # As usual, define some paths first to make life simpler
training_data = Path('./drive/My Drive/monkey species/training_data')
validation_data = Path('./drive/My Drive/monkey species/validation')
labels_path = Path('./drive/My Drive/monkey species/monkey_labels.txt')
```

```
[5] labels_info = []

# Read the file
lines = labels_path.read_text().strip().splitlines()[1:]
for line in lines:
    line = line.split(',')
    line = [x.strip('\n\t\r') for x in line]
    line[3], line[4] = int(line[3]), int(line[4])
    line = tuple(line)
    labels_info.append(line)

# Convert the data into a pandas dataframe
labels_info = pd.DataFrame(labels_info, columns=['Label', 'Latin Name', 'Common Name',
                                                'Train Images', 'Validation Images'], index=None)

# Sneak peek
labels_info.head(10)
```

| | Label | Latin Name | Common Name | Train Images | Validation Images |
|---|-------|----------------------|---------------------------|--------------|-------------------|
| 0 | n0 | alouatta_palliata | mantled_howler | 131 | 26 |
| 1 | n1 | erythrocebus_patas | patas_monkey | 139 | 28 |
| 2 | n2 | cacajao_calvus | bald_uakari | 137 | 27 |
| 3 | n3 | macaca_fuscata | japanese_macaque | 152 | 30 |
| 4 | n4 | cebuella_pygmea | pygmy_marmoset | 131 | 26 |
| 5 | n5 | cebus_capucinus | white_headed_capuchin | 141 | 28 |
| 6 | n6 | mico_argentatus | silvery_marmoset | 132 | 26 |
| 7 | n7 | saimiri_sciureus | common_squirrel_monkey | 142 | 28 |
| 8 | n8 | aotus_nigriceps | black_headed_night_monkey | 133 | 27 |
| 9 | n9 | trachypitecus_johnii | nilgiri_langur | 132 | 26 |

```
[8] # Create a dictionary to map the labels to integers
labels_dict= {'n0':0, 'n1':1, 'n2':2, 'n3':3, 'n4':4, 'n5':5, 'n6':6, 'n7':7, 'n8':8, 'n9':9}

# map labels to common names
names_dict = dict(zip(labels_dict.values(), labels_info["Common Name"]))
print(names_dict)
```

```
{0: 'mantled_howler', 1: 'patas_monkey', 2: 'bald_uakari', 3: 'japanese_macaque', 4: 'pygmy_marmoset', 5: 'white_headed_capuchin', 6: 'silvery_marmoset', 7: 'common_squirrel_monkey'}
```

```
# Creating a dataframe for the training dataset
train_df = []
for folder in os.listdir(training_data):
    # Define the path to the images
    imgs_path = training_data / folder

    # Get the list of all the images stored in that directory
    imgs = sorted(imgs_path.glob('*.jpg'))

    # Store each image path and corresponding label
    for img_name in imgs:
        train_df.append((str(img_name), labels_dict[folder]))

train_df = pd.DataFrame(train_df, columns=['image', 'label'], index=None)
# shuffle the dataset
train_df = train_df.sample(frac=1.).reset_index(drop=True)

#####

# Creating dataframe for validation data in a similar fashion
valid_df = []
for folder in os.listdir(validation_data):
    imgs_path = validation_data / folder
    imgs = sorted(imgs_path.glob('*.jpg'))
    for img_name in imgs:
        valid_df.append((str(img_name), labels_dict[folder]))

valid_df = pd.DataFrame(valid_df, columns=['image', 'label'], index=None)
# shuffle the dataset
valid_df = valid_df.sample(frac=1.).reset_index(drop=True)

#####

# How many samples do we have in our training and validation data?
print("Number of training samples: ", len(train_df))
print("Number of validation samples: ", len(valid_df))

# sneak peek of the training and validation dataframes
print("\n", train_df.head(), "\n")
print("=====\n")
print("\n", valid_df.head())
```

➞ Number of training samples: 1096
Number of validation samples: 272

| | image | label |
|---|---|-------|
| 0 | drive/My Drive/monkey spicies/training_data/n9... | 9 |
| 1 | drive/My Drive/monkey spicies/training_data/n6... | 6 |
| 2 | drive/My Drive/monkey spicies/training_data/n0... | 0 |
| 3 | drive/My Drive/monkey spicies/training_data/n6... | 6 |
| 4 | drive/My Drive/monkey spicies/training_data/n9... | 9 |

=====

| | image | label |
|---|---|-------|
| 0 | drive/My Drive/monkey spicies/validation/n8/n8... | 8 |
| 1 | drive/My Drive/monkey spicies/validation/n0/n0... | 0 |
| 2 | drive/My Drive/monkey spicies/validation/n8/n8... | 8 |
| 3 | drive/My Drive/monkey spicies/validation/n6/n6... | 6 |
| 4 | drive/My Drive/monkey spicies/validation/n6/n6... | 6 |

```
[10] # some constants(not truly though!)

# dimensions to consider for the images
img_rows, img_cols, img_channels = 224,224,3

# batch size for training
batch_size=8

# total number of classes in the dataset
nb_classes=10
```

```
[11] # Augmentation sequence
seq = iaa.OneOf([
    iaa.Fliplr(), # horizontal flips
    iaa.Affine(rotate=20), # roatation
    iaa.Multiply((1.2, 1.5))] #random brightness
```



```
[12] def data_generator(data, batch_size, is_validation_data=False):
    # Get total number of samples in the data
    n = len(data)
    nb_batches = int(np.ceil(n/batch_size))

    # Get a numpy array of all the indices of the input data
    indices = np.arange(n)

    # Define two numpy arrays for containing batch data and labels
    batch_data = np.zeros((batch_size, img_rows, img_cols, img_channels), dtype=np.float32)
    batch_labels = np.zeros((batch_size, nb_classes), dtype=np.float32)

    while True:
        if not is_validation_data:
            # shuffle indices for the training data
            np.random.shuffle(indices)

        for i in range(nb_batches):
            # get the next batch
            next_batch_indices = indices[i*batch_size:(i+1)*batch_size]

            # process the next batch
            for j, idx in enumerate(next_batch_indices):
                img = cv2.imread(data.iloc[idx]["image"])
                img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                label = data.iloc[idx]["label"]

                if not is_validation_data:
                    img = seq.augment_image(img)

                img = cv2.resize(img, (img_rows, img_cols)).astype(np.float32)
                batch_data[j] = img
                batch_labels[j] = to_categorical(label, num_classes=nb_classes)

            batch_data = preprocess_input(batch_data)
            yield batch_data, batch_labels
```

```
[13] #training data generator
train_data_gen = data_generator(train_df, batch_size)

# validation data generator
valid_data_gen = data_generator(valid_df, batch_size, is_validation_data=True)
```

```
[14] # simple function that returns the base model
def get_base_model():
    base_model = VGG16(input_shape=(img_rows, img_cols, img_channels), weights='imagenet', include_top=True)
    return base_model
```

```
[15] # get the base model
base_model = get_base_model()

# get the output of the second last dense layer
base_model_output = base_model.layers[-2].output

# add new layers
x = Dropout(0.7, name='drop2')(base_model_output)
output = Dense(10, activation='softmax', name='fc3')(x)

# define a new model
model = Model(base_model.input, output)

# Freeze all the base model layers
for layer in base_model.layers[:-1]:
    layer.trainable=False

# compile the model and check it
optimizer = RMSprop(0.001)
model.compile(loss='categorical_crossentropy', optimizer=optimizer, metrics=['accuracy'])
model.summary()
```


[15] Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

| Layer (type) | Output Shape | Param # |
|----------------------------|-----------------------|---------|
| ===== | | |
| input_1 (InputLayer) | (None, 224, 224, 3) | 0 |
| block1_conv1 (Conv2D) | (None, 224, 224, 64) | 1792 |
| block1_conv2 (Conv2D) | (None, 224, 224, 64) | 36928 |
| block1_pool (MaxPooling2D) | (None, 112, 112, 64) | 0 |
| block2_conv1 (Conv2D) | (None, 112, 112, 128) | 73856 |
| block2_conv2 (Conv2D) | (None, 112, 112, 128) | 147584 |
| block2_pool (MaxPooling2D) | (None, 56, 56, 128) | 0 |
| block3_conv1 (Conv2D) | (None, 56, 56, 256) | 295168 |
| block3_conv2 (Conv2D) | (None, 56, 56, 256) | 590080 |
| block3_conv3 (Conv2D) | (None, 56, 56, 256) | 590080 |
| block3_pool (MaxPooling2D) | (None, 28, 28, 256) | 0 |
| block4_conv1 (Conv2D) | (None, 28, 28, 512) | 1180160 |
| block4_conv2 (Conv2D) | (None, 28, 28, 512) | 2359808 |
| block4_conv3 (Conv2D) | (None, 28, 28, 512) | 2359808 |
| block4_pool (MaxPooling2D) | (None, 14, 14, 512) | 0 |

```
[16] # always user earlystopping
# the restore_best_weights parameter load the weights of the best iteration once the training finishes
es = EarlyStopping(patience=10, restore_best_weights=True)

# checkpoint to save model
chkpt = ModelCheckpoint(filepath="model1", save_best_only=True)

# number of training and validation steps for training and validation
nb_train_steps = int(np.ceil(len(train_df)/batch_size))
nb_valid_steps = int(np.ceil(len(valid_df)/batch_size))

# number of epochs
nb_epochs=1
```

```
[17] # train the model
history1 = model.fit_generator(train_data_gen,
                               epochs=nb_epochs,
                               steps_per_epoch=nb_train_steps,
                               validation_data=valid_data_gen,
                               validation_steps=nb_valid_steps,
                               callbacks=[es,chkpt])
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and w
Instructions for updating:
Use tf.cast instead.
Epoch 1/1
137/137 [=====] - 495s 4s/step - loss: 1.4765 - acc: 0.7016 - val_loss: 0.3190 - val_acc: 0.9338
```

```
[ ] # memory footprint support libraries/code
ln -sf /opt/bin/nvidia-smi /usr/bin/nvidia-smi
!pip install gputil
!pip install psutil
!pip install humanize
import psutil
import humanize
import os
import GPUtil as GPU
GPUs = GPU.getGPUs()
# XXX: only one GPU on Colab and isn't guaranteed
gpu = GPUs[0]
def printm():
    process = psutil.Process(os.getpid())
    print("Gen RAM Free: " + humanize.naturalsize( psutil.virtual_memory().available ), " | Proc size: " + humanize.naturalsize( process.memory_info().rss))
    print("GPU RAM Free: {0:.0f}MB | Used: {1:.0f}MB | Util {2:3.0f}% | Total {3:.0f}MB".format(gpu.memoryFree, gpu.memoryUsed, gpu.memoryUtil*100, gpu.memoryTotal))
    printm()
```

Collecting gputil
 Downloading <https://files.pythonhosted.org/packages/ed/0e/5c61eedde9f6c87713e89d794f01e378cf9565847d4576fa627d758c554/GPUtil-1.4.0.tar.gz>
 Building wheels for collected packages: gputil
 Building wheel for gputil (setup.py) ... done
 Stored in directory: /root/.cache/pip/wheels/3d/77/07/80562de4bb0786e5ea186911a2c831fdd0018bda69beab71fd
 Successfully built gputil
 Installing collected packages: gputil
 Successfully installed gputil-1.4.0
 Requirement already satisfied: psutil in /usr/local/lib/python3.6/dist-packages (5.4.8)
 Requirement already satisfied: humanize in /usr/local/lib/python3.6/dist-packages (0.5.1)
 Gen RAM Free: 9.3 GB | Proc size: 5.2 GB
 GPU RAM Free: 8384MB | Used: 6695MB | Util 44% | Total 15079MB

```
[ ] # get the training and validation accuracy from the history object
train_acc = history1.history['acc']
valid_acc = history1.history['val_acc']

# get the loss
train_loss = history1.history['loss']
valid_loss = history1.history['val_loss']

# get the number of entries
xvalues = np.arange(len(train_acc))

# visualize
f,ax = plt.subplots(1,2, figsize=(10,5))
ax[0].plot(xvalues, train_loss)
ax[0].plot(xvalues, valid_loss)
ax[0].set_title("Loss curve")
ax[0].set_xlabel("Epoch")
ax[0].set_ylabel("loss")
ax[0].legend(['train', 'validation'])

ax[1].plot(xvalues, train_acc)
ax[1].plot(xvalues, valid_acc)
ax[1].set_title("Accuracy")
ax[1].set_xlabel("Epoch")
ax[1].set_ylabel("accuracy")
ax[1].legend(['train', 'validation'])

plt.show()
```

```
# What is the final loss and accuracy on our validation data?
valid_loss, valid_acc = model.evaluate_generator(valid_data_gen, steps=nb_valid_steps)
print(f"Final validation accuracy: {valid_acc*100:.2f}%")
```

Final validation accuracy: 98.16%

Observations: We labeled species as n0, n1, ..., n9. And achieved validation accuracy as 98.5%.

Task4: Text classification using CNN:

For this we have used sentiment reviews dataset from Kaggle and implemented model and comments foreach line is explained below and had accuracy 98%

```
[1] from google.colab import drive
drive.mount('/content/drive')
```

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[2] import numpy as np
import pandas as pd
import nltk
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from bs4 import BeautifulSoup
import re
from keras.utils import to_categorical
import random
from tensorflow import set_random_seed
from sklearn.model_selection import train_test_split
from keras.preprocessing import sequence
from keras.preprocessing.text import Tokenizer
from keras.layers import Dense,Dropout,Embedding,LSTM
from keras.layers import Dense, Activation,Flatten
from keras.layers.convolutional import Conv1D,MaxPooling1D
from keras.callbacks import EarlyStopping
from keras.losses import categorical_crossentropy
from keras.optimizers import Adam
from keras.models import Sequential
from tqdm import tqdm
import warnings
warnings.filterwarnings("ignore", category=UserWarning, module='bs4')
lemmatizer = WordNetLemmatizer()

#set random seed for the session and also for tensorflow that runs in background for keras
set_random_seed(123)
random.seed(123)
```

↳ Using TensorFlow backend.

```
[3] train= pd.read_csv("./drive/My Drive/sentiment review/train.tsv", sep="\t")
test = pd.read_csv("./drive/My Drive/sentiment review/test.tsv", sep="\t")

train.head()
```

↳

| | PhraseId | SentenceId | Phrase | Sentiment |
|---|----------|------------|---|-----------|
| 0 | 1 | 1 | A series of escapades demonstrating the adage ... | 1 |
| 1 | 2 | 1 | A series of escapades demonstrating the adage ... | 2 |
| 2 | 3 | 1 | A series | 2 |
| 3 | 4 | 1 | A | 2 |
| 4 | 5 | 1 | series | 2 |

```
[4] train['Phrase'][200]
```

↳ ', Trouble Every Day is a plodding mess .'

```
[5] def clean_sentences(df):
    reviews = []

    for sent in tqdm(df['Phrase']):

        #remove html content
        review_text = BeautifulSoup(sent).get_text()

        #remove non-alphabetic characters
        review_text = re.sub("[^a-zA-Z]", " ", review_text)

        #tokenize the sentences
        words = word_tokenize(review_text.lower())

        #lemmatize each word to its lemma
        lemma_words = [lemmatizer.lemmatize(i) for i in words]

        reviews.append(lemma_words)

    return(reviews)
```

```
[6] import nltk
     nltk.download('all')
```

```
↳ [nltk_data] Downloading collection 'all'
[nltk_data] |
[nltk_data] | Downloading package abc to /root/nltk_data...
[nltk_data] | Package abc is already up-to-date!
[nltk_data] | Downloading package alpino to /root/nltk_data...
[nltk_data] | Package alpino is already up-to-date!
[nltk_data] | Downloading package biocreative_ppi to
[nltk_data] | /root/nltk_data...
[nltk_data] | Package biocreative_ppi is already up-to-date!
[nltk_data] | Downloading package brown to /root/nltk_data...
[nltk_data] | Package brown is already up-to-date!
↳ [nltk_data] | Downloading package mwa_ppdb to /root/nltk_data...
[nltk_data] | Package mwa_ppdb is already up-to-date!
[nltk_data] |
[nltk_data] Done downloading collection all
True
```

```
[7] #cleaned reviews for both train and test set retrieved
     train_sentences = clean_sentences(train)
     test_sentences = clean_sentences(test)
     print(len(train_sentences))
     print(len(test_sentences))
```

```
↳ 100%|██████████| 156060/156060 [01:02<00:00, 2486.25it/s]
100%|██████████| 66292/66292 [00:26<00:00, 2500.39it/s]156060
66292
```

```
[8] train_sentences[200]
```

```
↳ ['trouble', 'every', 'day', 'is', 'a', 'plodding', 'mess']
```

```
[9] target=train.Sentiment.values
     y_target=to_categorical(target)
     num_classes=y_target.shape[1]
```

```
[10] y_target
```

```
array([[0., 1., 0., 0., 0.],
       [0., 0., 1., 0., 0.],
       [0., 0., 1., 0., 0.],
       ...,
       [0., 0., 0., 1., 0.],
       [0., 0., 1., 0., 0.],
       [0., 0., 1., 0., 0.]], dtype=float32)
```

```
[11] X_train,X_val,y_train,y_val=train_test_split(train_sentences,y_target,test_size=0.2,stratify=y_target)
```

```
[12] #It is needed for initializing tokenizer of keras and subsequent padding
```

```
unique_words = set()
len_max = 0

for sent in tqdm(X_train):
    unique_words.update(sent)

    if(len_max<len(sent)):
        len_max = len(sent)

#length of the list of unique_words gives the no of unique words
print(len(list(unique_words)))
print(len_max)
```

```
100%|██████████| 124848/124848 [00:00<00:00, 528325.62it/s]13736
48
```

```
[13] for x in tqdm(X_train[1:10]):
    print(x)
```

```
100%|██████████| 9/9 [00:00<00:00, 4770.47it/s]['is', 'n', 't', 'necessarily']
['be', 'appreciated', 'by', 'anyone', 'outside', 'the', 'under', 'set']
['like', 'a', 'le', 'dizzily', 'gorgeous', 'companion', 'to', 'mr', 'wong', 's', 'in', 'the', 'mood', 'for', 'love', 'very', 'much', 'a', 'hong', 'kong', 'movie', 'despite', 'it',
['frailty', 'start', 'out', 'like', 'a', 'typical', 'bible', 'killer', 'story']
['retrieve']
['go', 'back', 'and', 'check', 'out', 'the', 'last', 'minute']
['glide', 'gracefully', 'from', 'male', 'persona', 'to', 'female']
['only', 'there', 'were', 'one', 'for', 'this', 'kind', 'of', 'movie']
['the', 'prospect', 'of', 'beck', 's', 'next', 'project']
```

```
[14] tokenizer = Tokenizer(num_words=len(list(unique_words)))
tokenizer.fit_on_texts(list(X_train))
X_train = tokenizer.texts_to_sequences(X_train)
X_val = tokenizer.texts_to_sequences(X_val)
X_test = tokenizer.texts_to_sequences(test_sentences)

#padding done to equalize the lengths of all input reviews. LSTM networks needs all inputs to be same length.
#Therefore reviews lesser than max length will be made equal using extra zeros at end. This is padding.
X_train = sequence.pad_sequences(X_train, maxlen=len_max)
X_val = sequence.pad_sequences(X_val, maxlen=len_max)
X_test = sequence.pad_sequences(X_test, maxlen=len_max)
print(X_train.shape,X_val.shape,X_test.shape)
```

```
(124848, 48) (31212, 48) (66292, 48)
```

```
[15] #Model using Keras CNN
model=Sequential()
model.add(Embedding(len(list(unique_words)),300,input_length=len_max))
model.add(Conv1D(128,5,activation='relu'))
model.add(MaxPooling1D(5))
model.add(Conv1D(128,5,activation='relu'))
model.add(MaxPooling1D(35))
model.add(Flatten())
model.add(Dense(100,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes,activation='softmax'))
model.compile(loss='categorical_crossentropy',optimizer=Adam(lr=0.005),metrics=['accuracy'])
model.summary()
```

```
[ ] history=model.fit(X_train, y_train, validation_data=(X_val, y_val),epochs=4, batch_size=256, verbose=1)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 124848 samples, validate on 31212 samples
Epoch 1/4
124848/124848 [=====] - 10s 81us/step - loss: 1.2020 - acc: 0.5367 - val_loss: 1.1307 - val_acc: 0.5650
Epoch 2/4
124848/124848 [=====] - 6s 49us/step - loss: 1.1045 - acc: 0.5808 - val_loss: 1.1028 - val_acc: 0.5778
Epoch 3/4
124848/124848 [=====] - 6s 48us/step - loss: 1.0632 - acc: 0.5961 - val_loss: 1.1013 - val_acc: 0.5796
Epoch 4/4
124848/124848 [=====] - 6s 47us/step - loss: 1.0364 - acc: 0.6047 - val_loss: 1.1014 - val_acc: 0.5818
```




```

import matplotlib.pyplot as plt

# Create count of the number of epochs
epoch_count = range(1, len(history.history['loss']) + 1)

# Visualize learning curve. Here learning curve is not ideal. It should be much smoother as it decreases.
#As mentioned before, altering different hyper parameters especially learning rate can have a positive impact
#on accuracy and learning curve.
plt.plot(epoch_count, history.history['loss'], 'r--')
plt.plot(epoch_count, history.history['val_loss'], 'b-')
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()

```



```

# What is the final loss and accuracy on our validation data?
valid_loss, valid_acc = model.evaluate_generator(valid_data_gen, steps=nb_valid_steps)
print(f"Final validation accuracy: {valid_acc*100:.2f}%")

```

Final validation accuracy: 98.16%

Task-5: Text classification Using LSTM model on sentiment reviews dataset

```

[ ]

[1] from google.colab import drive
drive.mount('/content/drive')

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client\_id=947318989803-6bn60k8qdgf4n4g3pfee6491hc0br4i.apps.googleusercontent.com&redirect\_uri=urn%3Aietf%3A
Enter your authorization code:
.....
Mounted at /content/drive

[2] import numpy as np
import pandas as pd
import nltk
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from bs4 import BeautifulSoup
import re
from keras.utils import to_categorical
import random
from tensorflow import set_random_seed
from sklearn.model_selection import train_test_split
from keras.preprocessing import sequence
from keras.preprocessing.text import Tokenizer
from keras.layers import Dense, Dropout, Embedding, LSTM
from keras.layers import Dense, Activation, Flatten
from keras.layers.convolutional import Conv1D, MaxPooling1D
from keras.callbacks import EarlyStopping
from keras.losses import categorical_crossentropy
from keras.optimizers import Adam
from keras.models import Sequential
from tqdm import tqdm
import warnings
warnings.filterwarnings("ignore", category=UserWarning, module='bs4')
lemmatizer = WordNetLemmatizer()

#set random seed for the session and also for tensorflow that runs in background for keras
set_random_seed(123)
random.seed(123)

Using TensorFlow backend.

```

```
[3] train= pd.read_csv("../drive/My Drive/sentiment review/train.tsv", sep="\t")
test = pd.read_csv("../drive/My Drive/sentiment review/test.tsv", sep="\t")
train.head()
```

```
PhraseId  SentenceId      Phrase  Sentiment
0         1          1  A series of escapades demonstrating the adage ...      1
1         2          1  A series of escapades demonstrating the adage ...      2
2         3          1                A series                          2
3         4          1                A                                2
4         5          1              series                          2
```

```
[7] train['Phrase'][200]
```

```
' , Trouble Every Day is a plodding mess . '
```

```
[8] def clean_sentences(df):
    reviews = []
    for sent in tqdm(df['Phrase']):
        #remove html content
        review_text = BeautifulSoup(sent).get_text()
        #remove non-alphabetic characters
        review_text = re.sub("[^a-zA-Z]", " ", review_text)
        #tokenize the sentences
        words = word_tokenize(review_text.lower())
        #lemmatize each word to its lemma
        lemma_words = [lemmatizer.lemmatize(i) for i in words]
        reviews.append(lemma_words)
    return(reviews)
```

```
[9] import nltk
nltk.download('all')
```

```
[nltk_data] Downloading collection 'all'
[nltk_data] |
[nltk_data] | Downloading package abc to /root/nltk_data...
[nltk_data] | Unzipping corpora/abc.zip.
[nltk_data] | Downloading package alpino to /root/nltk_data...
[nltk_data] | Unzipping corpora/alpino.zip.
[nltk_data] | Downloading package biocreative_ppi to
[nltk_data] | /root/nltk_data...
[nltk_data] | Unzipping corpora/biocreative_ppi.zip.
[nltk_data] | Downloading package brown to /root/nltk_data...
[nltk_data] | Unzipping corpora/brown.zip.
[nltk_data] | Downloading package brown_tei to /root/nltk_data...
```

```
[10] #cleaned reviews for both train and test set retrieved
train_sentences = clean_sentences(train)
test_sentences = clean_sentences(test)
print(len(train_sentences))
print(len(test_sentences))
```

```
100%|██████████| 156060/156060 [01:06<00:00, 2362.44it/s]
100%|██████████| 66292/66292 [00:27<00:00, 2412.35it/s]
156060
66292
```

```
[11] train_sentences[200]
```

```
['trouble', 'every', 'day', 'is', 'a', 'plodding', 'mess']
```

```
[12] target=train.Sentiment.values
y_target=to_categorical(target)
num_classes=y_target.shape[1]
```

```
[13] y_target
```

```
array([[0., 1., 0., 0., 0.],
       [0., 0., 1., 0., 0.],
       [0., 0., 1., 0., 0.],
       ...,
       [0., 0., 0., 1., 0.],
       [0., 0., 1., 0., 0.],
       [0., 0., 1., 0., 0.]], dtype=float32)
```



```
[14] X_train,X_val,y_train,y_val=train_test_split(train_sentences,y_target,test_size=0.2,stratify=y_target)
```

```
[15] #It is needed for initializing tokenizer of keras and subsequent padding
```

```
unique_words = set()
len_max = 0

for sent in tqdm(X_train):
    unique_words.update(sent)
    if(len_max<len(sent)):
        len_max = len(sent)

#length of the list of unique_words gives the no of unique words
print(len(list(unique_words)))
print(len_max)
```

```
100%|██████████| 124848/124848 [00:00<00:00, 463716.09it/s]13733
48
```

```
[16] for x in tqdm(X_train[1:10]):
    print(x)
```

```
100%|██████████| 9/9 [00:00<00:00, 3186.89it/s]['i', 'miss', 'something']
['on', 'dvd']
['s', 'in', 'the', 'mood', 'for', 'love', 'very', 'much', 'a', 'hong', 'kong', 'movie']
['you', 'would', 'n', 't', 'want', 'to', 'live', 'waydowntown', 'but', 'it', 'is', 'a', 'hilarious', 'place', 'to', 'visit']
['in', 'this', 'summer', 's', 'new', 'action', 'film']
['not', 'one', 'moment', 'in', 'the', 'enterprise']
['wide', 'screen']
['dialogue', 'rip']
['will', 'leave', 'the', 'auditorium', 'feeling', 'dizzy', 'confused', 'and', 'totally', 'disorientated']
```

```
[17] tokenizer = Tokenizer(num_words=len(list(unique_words)))
tokenizer.fit_on_texts(list(X_train))
X_train = tokenizer.texts_to_sequences(X_train)
X_val = tokenizer.texts_to_sequences(X_val)
X_test = tokenizer.texts_to_sequences(test_sentences)

#padding done to equalize the lengths of all input reviews. LSTM networks needs all inputs to be same length.
#Therefore reviews lesser than max length will be made equal using extra zeros at end. This is padding.
X_train = sequence.pad_sequences(X_train, maxlen=len_max)
X_val = sequence.pad_sequences(X_val, maxlen=len_max)
X_test = sequence.pad_sequences(X_test, maxlen=len_max)
print(X_train.shape,X_val.shape,X_test.shape)
```

```
(124848, 48) (31212, 48) (66292, 48)
```

```
[18] early_stopping = EarlyStopping(min_delta = 0.001, mode = 'max', monitor='val_acc', patience = 2)
callback = [early_stopping]
```

```
#Model using Keras LSTM
model=Sequential()
model.add(Embedding(len(list(unique_words)),300,input_length=len_max))
model.add(LSTM(128,dropout=0.5, recurrent_dropout=0.5,return_sequences=True))
model.add(LSTM(64,dropout=0.5, recurrent_dropout=0.5,return_sequences=False))
model.add(Dense(100,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes,activation='softmax'))
model.compile(loss='categorical_crossentropy',optimizer=Adam(lr=0.005),metrics=['accuracy'])
model.summary()
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated
Instructions for updating:
Colocations handled automatically by placer.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is depre
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
```

| Layer (type) | Output Shape | Param # |
|-------------------------|-----------------|---------|
| ***** | | |
| embedding_1 (Embedding) | (None, 48, 300) | 4119900 |

```
[18]
┌──────────┬──────────┬──────────┐
└lstm_1 (LSTM)          (None, 48, 128)      219648
└lstm_2 (LSTM)          (None, 64)          49408
└dense_1 (Dense)        (None, 100)         6500
└dropout_1 (Dropout)    (None, 100)          0
└dense_2 (Dense)        (None, 5)           505
└──────────┴──────────┴──────────┘
Total params: 4,395,961
Trainable params: 4,395,961
Non-trainable params: 0
```

```
[ ] !rm -R ./logs/ # rf
```

```
rm: cannot remove './logs/': No such file or directory
```


```
[19] history=model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=4, batch_size=256, verbose=1)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated. Instructions for updating: Use tf.cast instead.
Train on 124848 samples, validate on 31212 samples
Epoch 1/4
124848/124848 [=====] - 95s 760us/step - loss: 1.0053 - acc: 0.5954 - val_loss: 0.8407 - val_acc: 0.6542
Epoch 2/4
124848/124848 [=====] - 92s 738us/step - loss: 0.8056 - acc: 0.6703 - val_loss: 0.8219 - val_acc: 0.6670
Epoch 3/4
124848/124848 [=====] - 92s 741us/step - loss: 0.7356 - acc: 0.6935 - val_loss: 0.8118 - val_acc: 0.6741
Epoch 4/4
124848/124848 [=====] - 90s 725us/step - loss: 0.6944 - acc: 0.7083 - val_loss: 0.8313 - val_acc: 0.6708
```

```
import matplotlib.pyplot as plt

# Create count of the number of epochs
epoch_count = range(1, len(history.history['loss']) + 1)

# Visualize learning curve. Here learning curve is not ideal. It should be much smoother as it decreases.
# As mentioned before, altering different hyper parameters especially learning rate can have a positive impact
# on accuracy and learning curve.
plt.plot(epoch_count, history.history['loss'], 'r--')
plt.plot(epoch_count, history.history['val_loss'], 'b-')
plt.legend(['Training Loss', 'Validation Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()
```



```
[ ] # What is the final loss and accuracy on our validation data?
valid_loss, valid_acc = model.evaluate_generator(valid_data_gen, steps=nb_valid_steps)
print(f"Final validation accuracy: {valid_acc*100:.2f}%")
```

Final validation accuracy: 95.96%

Task-6: Comparison of CNN and LSTM on text classification for sentiment reviews data

As above outputs are showing that CNN has 98% accuracy where as LSTM has 95% accuracy on same dataset for the epochs size four. Hence CNN has right hand compared to LSTM slightly but not major change .Hyper parameter tuning for attaining above result have shown in above code snippets.

References:

<https://towardsdatascience.com/>
<https://www.kaggle.com/datasets>